

Principal Component Analysis in R

STATS 503


Example: Athletic Performance Data

We start reading the file `athletic.txt`. The first two columns correspond to the code and name of a country. The next columns correspond to the records for each competition.

```
Event <- c("X100", "X200", "X400", "X800", "X1500", "X5000", "X10000", "Marathon")
athletic_data <- read.table("athletic.txt", col.names=c("Code", "Country", Event))
Code <- athletic_data$Code
Country <- athletic_data$Country
numerical.data <- athletic_data[, 3:ncol(athletic_data)]
```

To produce pairwise scatterplots and histograms, we use the function `ggpairs` from the library `GGally`.

```
library(GGally)
ggpairs(numerical.data, axisLabels="none", diag=list(continuous=wrap('barDiag', bins=15)),
  upper=list(continuous="points"),
  lower=list(continuous="points"))
```




The function `princomp` returns object with principal component analysis of the data set. The argument `cor` decides whether it works with correlation or covariance matrix.

```
athletic.PCAcov <- princomp(numerical.data, cor=F)
loadings(athletic.PCAcov)[,1:2]
```

##	Comp.1	Comp.2
## X100	-0.019865414	-0.21068624
## X200	-0.041566107	-0.35895599
## X400	-0.110631839	-0.82784910
## X800	-0.005487697	-0.02317370
## X1500	-0.014386824	-0.04465503
## X5000	-0.079308429	-0.12996775
## X10000	-0.181098930	-0.29885158
## Marathon	-0.972786963	0.18081152

To plot a projection of the athletic events on the space spanned by the first two principal components, we can use the function `qplot`.

```
library(ggplot2)
athletic.PCcov <- predict(athletic.PCAcov)
qplot(x=athletic.PCcov[,1], y=athletic.PCcov[,2], label=Code, geom="text")
```




Now, we do the same way with the correlation matrix.

```
athletic.PCAcor <- princomp(numerical.data, cor=T)
loadings(athletic.PCAcor)[,1:2]

##           Comp.1      Comp.2
## X100    -0.3175502 -0.56692566
## X200    -0.3369774 -0.46154846
```

```
## X400      -0.3556427 -0.24831328
## X800      -0.3686786 -0.01239489
## X1500     -0.3728159  0.13973247
## X5000     -0.3643787  0.31201256
## X10000    -0.3667744  0.30686895
## Marathon -0.3419290  0.43898707
```


```
athletic.PCcor <- predict(athletic.PCAcor)
qplot(x=athletic.PCcor[,1], y=athletic.PCcor[,2], label=Code, geom="text")
```



Finally, the next code shows how to construct a barplot of the principal components' variance.

```
barplot(athletic.PCAcor$sdev^2)
title("Scree diagram from correlations")
```

Scree diagram from correlations



An example in high dimensions


The next code simulates data from a multivariate normal distribution with mean zero and covariance $\Sigma = \text{diag}(3^2, 2^2, 1, \dots, 1)$.

```
scree_plot <- function(n) {  
  p <- 100  
  X <- sweep(matrix(rnorm(n*p), ncol = p), 2, c(3, 2, rep(1, p-2)), "*")  
  W <- prcomp(X)  
  qqplot(y=W$sdev, xlab="index", ylab="Standard deviations", main=paste("n =", n),  
         ylim = c(-0.3, 3.5))  
}
```

```
plotlist = lapply(c(1000,100,20),screes_plot)

library(gridExtra)
grid.arrange(plotlist[[1]],plotlist[[2]],plotlist[[3]],ncol=3)


## Warning: Removed 1 rows containing missing values (geom_point).
```



Example: City crime data

```
data_crime <- read.table("citycrime.txt", header = TRUE)
ggpairs(data_crime,axisLabels="none", diag=list(continuous=wrap('barDiag', bins=15)),
  upper=list(continuous = "points"),
```

```
lower=list(continuous = "points"))
```



```
city.PCA <- princomp(data_crime,cor=T)
loadings(city.PCA)[,1:3]
```

```
##          Comp.1      Comp.2      Comp.3
## Murder  -0.3703162 -0.3393305  0.20197401
## Rape    -0.2494347  0.4665068  0.78285390
## Robbery  -0.4260252 -0.3878604  0.07906463
## Assault  -0.4340165  0.0424326 -0.28183161
## Burglary -0.4497241  0.2382532  0.01503743
## Larceny  -0.2759217  0.6055442 -0.49241798
## MVT      -0.3903791 -0.3025585 -0.13403135
```


```
summary(city.PCA)
```

```
## Importance of components:
```

```
##           Comp.1    Comp.2    Comp.3    Comp.4    Comp.5
## Standard deviation  1.948024 1.0950164 0.8771129 0.71434047 0.57690506
## Proportion of Variance 0.542114 0.1712944 0.1099039 0.07289747 0.04754564
## Cumulative Proportion 0.542114 0.7134084 0.8233123 0.89620976 0.94375539
##           Comp.6    Comp.7
## Standard deviation  0.4617666 0.42483396
## Proportion of Variance 0.0304612 0.02578341
## Cumulative Proportion 0.9742166 1.00000000
```

```
cityPC <- predict(city.PCA)
```

```
qplot(x=cityPC[,1], y=cityPC[,2], label=row.names(data_crime), geom="text")
```

The next code generates 10,000 data sets X^* taken from the dataset `data_crime` with replacement. Then, it calculates the percentage of the variance explained by the first two principal components.

```
N <- 10000
M <- length(data_crime[[1]])
data_crime = as.data.frame(data_crime)
theta_samples <- sapply(1:N, function(i){
  X_star <- data_crime[sample(1:length(data_crime[[1]]),size=M,replace=TRUE),]
  u <- summary(princomp(X_star,cor = T))
  return(sum(u[[1]][1:2]^2)/sum(u[[1]]^2))
})


sample_percentage <- sum(city.PCA$sdev[1:2]^2)/(sum(city.PCA$sdev^2))

qplot(theta_samples, geom="histogram",bins=100, fill=I("blue"),
```

```

xlab="Percentage of variance") +
geom_vline(aes(xintercept=sample_percentage,color="Sample.percentage"), show.legend = TRUE)+
scale_colour_manual(name="", values=c(Sample.percentage="red"))

```



```

quantile(theta_samples,probs = 0.025)

##      2.5%
## 0.6633778

quantile(theta_samples,probs = 0.975)

##      97.5%
## 0.7775128

```


Now, we calculate the variance explained by components 1 and 2 in 10,000 generated data sets.

```


lambdas <- sapply(1:N, function(i){
  X_permuted <- sapply(data_crime, function(x){return(x[sample(M)])})
  pca <- princomp(X_permuted, cor=T)
  return(pca$sdev^2)})
eig1 <- city.PCA$sdev[1]^2
eig2 <- city.PCA$sdev[2]^2
qplot(lambdas[1,], geom="histogram", bins=100, fill=I("blue"), xlab="First eigenvalue",
      xlim = c(1,4)) + geom_vline(aes(xintercept=eig1, color="Eig1"), show.legend = TRUE) +
  scale_colour_manual(name="", values=c(Eig1="red"))

## Warning: Removed 3 rows containing missing values (geom_bar).

```




```
qplot(lambdas[2,], geom="histogram",bins=100, fill=I("blue"),xlab="Second eigenvalue") + geom_vline(aes(
  scale_colour_manual(name="", values=c(Eig2="red"))
```



Toy example

```
p <- 2
Sigma <- matrix(c(2, 1, 1, 1), nrow=2)
n <- 30
X <- data.frame(t(chol(Sigma)%*%matrix(rnorm(p*n),nrow=p)))
pca <- prcomp(X)


qplot(x=X[,1], y=X[,2], label=as.character(1:n), geom="text", main="Data")
```



```
#library(devtools)
#install_github("vqv/ggbiplot")
library(ggbiplot)

## Loading required package: plyr
## Loading required package: scales
## Loading required package: grid

ggbiplot(pcoobj=pca, labels=as.character(1:n)) + ggtitle("Biplot")
```



Example: Applicant's qualifications

```
applicants <- read.table("applicants_qualifications.txt", header=TRUE)
applicantsPCA <- princomp(applicants, cor=T)
loadings(applicantsPCA)[,1:4]
```

##	Comp.1	Comp.2	Comp.3	Comp.4
## Formof.letter	-0.16244045	0.42884630	-0.31537452	0.09434720
## Appearance	-0.21310790	-0.03526587	0.02287769	-0.26217502
## Academic	-0.04018392	0.23691912	0.43046955	-0.63627435
## Likeability	-0.22507847	-0.12979571	-0.46582469	-0.34537465
## Self.conf	-0.29048107	-0.24889621	0.24102627	0.17280371


```
## Lucidity      -0.31487036 -0.13099029  0.15003717  0.07103303
## Honesty       -0.15811663 -0.40544998 -0.28392797 -0.41649052
## Salesmanship  -0.32425606 -0.02949230  0.18597471  0.19822750
## Experience    -0.13406824  0.55313856 -0.08259060 -0.06775209
## Drive         -0.31507146  0.04624321  0.07963541  0.15598748
## Ambition      -0.31802398 -0.06815490  0.20865130  0.19929112
## Grasp         -0.33149687 -0.02315034  0.11714220 -0.07472588
## Potential     -0.33328897  0.02225729  0.07254382 -0.18814004
## Keen.to.join  -0.25920844 -0.08227158 -0.46720556  0.20137601
## Suitability   -0.23603667  0.42066207 -0.08915180  0.01991326
```

```
summary(applicantsPCA)
```

```
## Importance of components:
```


```
##              Comp.1    Comp.2    Comp.3    Comp.4
## Standard deviation  2.7411301 1.4339809 1.20657345 1.09448513
## Proportion of Variance 0.5009196 0.1370867 0.09705463 0.07985985
## Cumulative Proportion 0.5009196 0.6380064 0.73506099 0.81492084
##              Comp.5    Comp.6    Comp.7    Comp.8
## Standard deviation  0.85973985 0.70326316 0.59267346 0.55668844
## Proportion of Variance 0.04927684 0.03297194 0.02341746 0.02066013
## Cumulative Proportion 0.86419768 0.89716961 0.92058707 0.94124720
##              Comp.9    Comp.10    Comp.11    Comp.12
## Standard deviation  0.50691374 0.43001206 0.39074335 0.312350893
## Proportion of Variance 0.01713077 0.01232736 0.01017869 0.006504205
## Cumulative Proportion 0.95837797 0.97070533 0.98088402 0.987388228
##              Comp.13    Comp.14    Comp.15
## Standard deviation  0.298024834 0.254230665 0.189009390
## Proportion of Variance 0.005921253 0.004308882 0.002381637
## Cumulative Proportion 0.993309481 0.997618363 1.000000000
```

```
ggbiplot(applicantsPCA)
```



Example: Canadian temperatures

```
canadian_weather <- data.frame(t(read.table("canadian_weather.txt", header=TRUE)))
canadian_weather$month = factor(row.names(canadian_weather))
library(reshape2)
meltedf <- melt(canadian_weather, "month")
meltedf$month = factor(meltedf$month, levels=row.names(canadian_weather))
ggplot(meltedf, aes(x=month, y=value, group=variable, color=variable)) + geom_line()
```

```
weatherPCA <- princomp(t(canadian_weather[,1:35]), cor=T)
summary(weatherPCA)


## Importance of components:
##               Comp.1    Comp.2    Comp.3    Comp.4
## Standard deviation  3.195162  1.2076746  0.47997719  0.261666684
## Proportion of Variance 0.850755  0.1215398  0.01919818  0.005705788
## Cumulative Proportion 0.850755  0.9722949  0.99149305  0.997198833
##               Comp.5    Comp.6    Comp.7    Comp.8
## Standard deviation  0.134133729  0.0876024923  0.0629457104  0.0402357196
## Proportion of Variance 0.001499321  0.0006395164  0.0003301802  0.0001349094
## Cumulative Proportion 0.998698154  0.9993376709  0.9996678511  0.9998027605
##               Comp.9    Comp.10    Comp.11    Comp.12
## Standard deviation  3.227973e-02  2.744603e-02  1.877478e-02  1.480256e-02
```

```
## Proportion of Variance 8.683175e-05 6.277372e-05 2.937436e-05 1.825966e-05
## Cumulative Proportion 9.998896e-01 9.999524e-01 9.999817e-01 1.000000e+00
```

```
loadings(weatherPCA)[,1:3]
```

```
##          Comp.1      Comp.2      Comp.3
## Jan -0.2742230 -0.38903639  0.00253309
## Feb -0.2851561 -0.31426293 -0.29287615
## Mar -0.3023976 -0.15439433 -0.35285691
## Apr -0.3036991  0.06946101 -0.42350861
## May -0.2912423  0.25922541 -0.37012616
## Jun -0.2658215  0.42435087 -0.13737141
## Jul -0.2614086  0.44005199  0.22378267
## Aug -0.2884328  0.29264101  0.29102415
## Sep -0.3085248  0.07180214  0.27528430
## Oct -0.3069895 -0.05898804  0.24294377
## Nov -0.2918766 -0.23661894  0.40088871
## Dec -0.2796017 -0.36062087  0.15555290
```

```
weatherPC <- predict(weatherPCA)
qplot(x=weatherPC[,1], y=weatherPC[,2], label=colnames(canadian_weather)[1:35],
      geom = "text")
```



```
ggbiplot(weatherPCA)
```

